

Frequency or Information? A Diagnostic Study of Spectral Side-Channel Conditioning in Graph Diffusion

Tamara Bluzer Dan Shamia Daniel Halperin Itamar Kolodny

315287441 208004119 207826314 211490362

Tel Aviv University

Abstract

We test whether spectral input helps graph generation or mainly reveals the source graph. Our diffusion model generates adjacency matrices with Laplacian eigenpairs as extra input. On Planar, five conditioning choices and five cutoffs give two strong settings: low-frequency $k=8$ and high-frequency $k=32$ reduce MMD Ratio by factors of 2.3 and 29.7 across three seeds. We use *oracle conditioning*: input from real validation graphs, not independently generated spectra. A simple reconstruction test with access to target details recovers 72.7% and 98.5% of the targets' edges. Across 12 saved pilot runs, low and high at $k=8$ each identify their source graph uniquely in all 96 samples. Recovery does not always track quality, but this does not rule out copying. Classification results remain inconclusive. Reconstruction should therefore be checked before attributing generation gains to frequency content.

1 Introduction

We generate graph connections without node or edge attributes. An adjacency matrix records which nodes are connected. Diffusion models learn to remove noise and generate graphs starting from noise (Niu et al., 2020; Jo et al., 2022; Vignac et al., 2023). Laplacian eigenpairs give extra structural information (Martinkus et al., 2022; Minello et al., 2025). Each contains an eigenvalue and its eigenvector. The *band* selects which eigenpairs we supply; the *cutoff* k is how many.

Motivated by SDMG (Zhu et al., 2025), our proposal asked whether low-frequency input improves latent diffusion and whether adding more eigenpairs eventually hurts. We keep this question but generate adjacency directly, not a learned latent representation. The eigenpairs form a *side channel*: extra input that guides denoising. Changing k changes this input, not the graph representation. But different bands can reveal different

amounts of the source graph. We ask: *which inputs improve scores, and how much do the outputs reconstruct the graphs that supplied them?*

We assess generation scores alongside edge recovery from the condition and overlap with source graphs. These checks reveal reconstruction, not a causal frequency effect. Code, results and analyzed samples accompany this paper.¹

2 Related Work

SPECTRE (Martinkus et al., 2022) generates eigenpairs with a GAN and decodes an adjacency matrix; we reuse its cutoff grid. GGSD (Minello et al., 2025) already compares diffusion representations from the smallest and largest eigenpairs. We instead use eigenpairs as extra input and measure reconstruction. DiGress (Vignac et al., 2023) supplies our backbone and benchmark conventions, but our main process is continuous, not discrete. LGD (Zhou et al., 2024) motivated our latent approach. DualDiff (Xie and Pan, 2026) combines node and cluster latents, rather than frequency-ordered inputs. SDMG (Zhu et al., 2025) studies representation learning; its findings need not transfer to generation.

3 Method and Scope

Data and task. We use SPECTRE's 200 Planar graphs: Delaunay triangulations on 64 nodes. Following DiGress, seed 0 splits them into 128 training, 32 validation and 40 test graphs without duplicates. Training teaches the generator; validation selects its checkpoint. Each run generates 32 graphs. **We also score generation on validation**, not the untouched test set. Classification uses separate synthetic splits (§4.6).

What each setting receives. A *condition* is extra input to the denoiser. For adjacency A and de-

¹https://github.com/TamaraBluzer/Frequency_aware_diffusion-

gree matrix D , we compute the normalized Laplacian $L = I - D^{-1/2}AD^{-1/2} = U\Lambda U^\top$. Eigenvalues are sorted from smallest to largest; n is the node count and c counts zero eigenvalues. For $k \in \{2, 4, 8, 16, 32\}$, we compare:

low and high: the first k nonzero eigenpairs (indices $c + 1, \dots, c + k$) or the last k ($n - k + 1, \dots, n$). This tests which end of the spectrum helps: smoother patterns or faster variation between neighboring nodes.

random: k randomly selected nonzero eigenpairs of the *same graph*. They still carry real structure, but do not favor either frequency band.

gaussian: independent $\mathcal{N}(0, 1)$ numbers replace the k eigenvalues and $n \times k$ eigenvector entries. This tests extra input without target-edge information. It is noise in the hint, not the diffusion noise used by every model.

shuffled: each graph receives a *different graph's* low-frequency condition. We shuffle graph-condition assignments, not node order, with no self-matches. This tests whether real spectral structure helps without the correct target pairing.

none: zero-valued condition inputs and an “absent” flag give no spectral hint. The model still receives the noisy adjacency, timestep and graph size. This is our baseline for measuring the benefit of conditioning.

Conditions stay fixed while denoising. Training uses training-graph conditions; oracle sampling uses validation-graph conditions. To supply node-pair information, we form $P_k = U_k \text{diag}(\lambda_k) U_k^\top$ and scale it to unit root-mean-square magnitude, so bands do not differ simply in input scale. Real eigenvalues are mapped to $[-1, 1]$ and supplied globally. Gaussian substitutes use the same processing. P_k is unchanged by sign flips or rotations within a fully selected repeated eigenspace, but not necessarily when that eigenspace is split (Lim et al., 2023). Parameter counts match at fixed k ; larger k widens the global input projection.

Model. We corrupt $x_0 = 2A - 1$ with symmetric Gaussian noise under a cosine schedule and train the model to predict the clean matrix (Ho et al., 2020; Nichol and Dhariwal, 2021). Sampling starts from noise; after denoising, positive values become edges. The transformer is permutation-equivariant: renaming nodes correspondingly renames its outputs. It has four layers, $(d_x, d_e, d_y) = (128, 64, 128)$, eight heads and 1.83M parameters at $k=8$. Training uses 200

epochs, batch 16, learning rate 3×10^{-4} and 100 diffusion steps. Dropping conditions with probability 0.1 enables classifier-free guidance, which adjusts reliance on the hint during sampling (Ho and Salimans, 2022). Self-loops and padding are excluded from the generated graph.

Changes from the proposal. We kept the Planar comparisons and added controls and reconstruction checks. Our autoencoder used 32,768 floats for 2,016 edge decisions and reconstructed perfectly by copying edges into latent entries. We switched to adjacency diffusion; this does not invalidate latent diffusion generally. Latent transfer and SBM generation remain untested; SBM results cover recovery only. We replaced planar-versus-SBM classification with community-count prediction to reduce dataset-identity shortcuts. Its high-frequency inputs and unverified labels prevent testing the proposed low-frequency benefit.

How we measure quality. Maximum mean discrepancy (MMD) (Gretton et al., 2012) compares statistics across real and generated graph sets, rather than demanding exact copies. We measure degree (neighbor counts), clustering (triangles), Laplacian spectrum, 12 ab-spline wavelet scales (multiscale structure) and four-node orbit counts (small patterns) (You et al., 2018; Liao et al., 2019; Hočevar and Demšar, 2014). Total-variation Gaussian kernel bandwidths are 1, 0.1, 1, 1, 30, respectively. Each MMD is divided by its train-validation real-data reference; a 10^{-8} denominator floor avoids unstable division. **Ratio** averages these relative scores; lower is better. A score of 1 matches the reference but does not establish statistical indistinguishability. Separately, validity checks that a graph is connected and can be drawn without edge crossings (planar) (Hagberg et al., 2008).

Random seeds change training and sampling choices. The pilot and selected-cutoff repeats use three seeds per setting; the full sweep and discrete comparisons use one; classifiers use five. Sample standard deviations (SD, ddof = 1) show variation across seeds. Pilot bootstrap intervals resample both graphs and seeds.

How we check reconstruction. To test how much the hint already reveals, our model-free *probe* ranks entries of P_k for distinct node pairs and keeps the target’s true edge count. It tries highest-first and lowest-first rankings, choosing

band	Ratio ↓	val. loss ↓	valid
none	327.2 ± 8.8	.132 ± .012	0%
low	158.0 ± 16.2	.053 ± .004	0%
high	339.1 ± 9.2	.094 ± .009	0%
random	326.5 ± 11.8	.128 ± .012	0%

Table 1: Oracle Planar pilot at $k=8$: validation scores, three-seed mean $\pm$ sample SD. Five-metric Ratio and loss both favor low, but rank the other settings differently.

component	Ratio ↓	none	low, $k=8$
Degree	138.3 ± 62.3	199.9 ± 11.1	
Clustering	17.9 ± 0.3	14.9 ± 0.0	
Spectrum	13.1 ± 0.7	7.3 ± 0.3	
Wavelet	356.4 ± 2.7	128.9 ± 1.6	
Orbit	1110.5 ± 21.9	438.8 ± 85.4	

Table 2: Individual pilot metrics, three-seed mean $\pm$ sample SD. Low improves four metrics but worsens degree. The 0.0 SD is a rounded value.

the one recovering more true edges. These target details make the test *privileged*; it is not independent generation. Low recovery does not rule out a better decoder. Gaussian controls measure the rule’s noise level, and the full spectrum gives complete recovery.

4 Results

4.1 Initial comparison at $k = 8$

Low-band input roughly halves Ratio and reduces validation loss by 60% (Table 1). The saved hierarchical bootstrap analysis (1,000 resamples) gives 95% intervals for low minus none, high and random of $[-189.4, -152.1]$, $[-206.7, -152.4]$ and $[-195.1, -142.2]$. With diffusion noise held fixed, adding the low condition changes 12.95% of edge decisions.

Most of the overall improvement comes from orbit counts, while degree gets worse (Table 2). A lower Ratio does not mean every structural metric improves. We have these individual metrics for the pilot, but not for the cutoff sweep.

4.2 Changing the cutoff

More eigenpairs do not always help (Figure 1). In seed 0, low $k=8$ scores 116.3 and high $k=32$ scores 10.5, versus 335.7 for none. Other scores range from 270.5 to 347.9; smaller gains have not been repeated. Gaussian ranges from 322.1 to 332.5 and shuffled from 328.8 to 342.0. These ranges do not measure baseline variation across

arm	Ratio ↓	base/arm	valid
none	330.1 ± 7.5	—	0%
low, $k=8$	141.9 ± 23.6	2.3	0%
high, $k=32$	11.1 ± 0.7	29.7	4.2%

Table 3: Repeated settings: three-seed mean $\pm$ sample SD, five-metric validation Ratio. Seed 0 also appears in the sweep. This is a separate set of runs from Table 1: low- $k=8$ averages 141.9 here and 158.0 there. High has 1/32, 2/32 and 1/32 valid samples across seeds.

band	$k=2$	4	8	16	32
low	48.1	54.0	72.7	78.7	47.9
high	21.8	31.4	49.3	76.7	98.5
random	14.8	15.6	19.2	26.3	43.6
gaussian	9.2	8.5	9.2	8.7	9.4
shuffled	8.7	8.7	8.8	8.7	8.3

Table 4: Edges recovered by the privileged probe (%) on Planar validation graphs. Bold marks the selected generation settings. Random-edge chance is 8.8%; Gaussian controls also account for choosing the better score orientation.

seeds.

Seeds 1 and 2 confirm gains, not optimal cut-offs.

4.3 How much of the target can be recovered?

Both selected conditions reveal many target edges (Table 4). At $k=8$, low has both the best generation score and the highest recovery. High $k=32$ allows 98.5% recovery, nearly reconstructing the target from its condition alone.

Reading the two axes. Low- $k=8$ and high- $k=16$ have similar recovery (72.7% and 76.7%) but different seed-0 Ratios (116.3 and 281.2) in Figure 1(b). At $k=16$, low and high have both similar recovery (78.7% and 76.7%) and similar Ratios (270.5 and 281.2). These are single-seed scores, not the low- $k=8$ three-seed mean of 141.9; neither $k=16$ setting was repeated. Recovery alone therefore does not consistently rank quality. Still, similar scores from one probe do not mean the inputs contain equal information. The first pair also differs in k and graph-level input width. High- $k=32$ shows how much the hint reveals; probe recovery does not measure copying by the generator.

Low recovery falls to 47.9% at $k=32$, while high recovery increases. A model need not use all the information that the probe can recover. On SBM, low recovery rises from 33.8% to 74.4% and




Figure 1: Oracle Planar sweep, five-metric validation Ratio. Hollow markers show seed 0; stars show three-seed means and sample SD for repeated settings. The dashed baseline and gray band show the unconditioned mean and SD, not significance. (a) Quality versus cutoff; lines guide the eye. (b) Quality versus probe recovery.

band	matched F_1	mismatched F_1	top-1
none	.085 $\pm$.004	.085	2/96
low	.617 $\pm$.005	.088	96/96
high	.337 $\pm$.005	.087	96/96
random	.108 $\pm$.003	.086	18/96

Table 5: Copying check, 12 saved $k=8$ runs. Matched: source graph, three-seed mean $\pm$ sample SD; mismatched: other 31 graphs. Top-1 counts unique nearest matches; chance is 1/32.

high from 23.0% to 87.7%, above 9.4% chance. This supports the recovery finding on another dataset, not the generation finding.

Using stored node order, we compare sample i with source graph i . Low and high identify their source uniquely in all 96 samples per band (Table 5). None has almost equal matched and mismatched overlap; random has a smaller gap than low or high. Figure 2 shows one example.

The snapshot records source hashes; older analysis used different samples. Original MMD reports lack sample hashes, preventing exact linkage. The copying analysis in Table 5 covers 12 runs at $k=8$.

4.4 Few generated graphs are valid

All 384 saved $k=8$ samples are nonplanar, although 67.7% of low and 91.7% of unconditioned samples are connected. High- $k=32$ repeat reports include 4/96 valid graphs. Separately, its probe gives 1/32 planar graphs despite 98.5% recovery.

Near-complete edge recovery does not guarantee planarity with this decoder; it does not follow that validity requires exact copying or that the dataset causes the failures.

4.5 Using more compute

In a separate five-metric discrete model (Austin et al., 2021), ten layers, 500 steps and 1,000 epochs improve unconditioned Ratio from 306.78 to 39.60. Conditioning reaches 15.40, shrinking the gap from 276.7 to 24.2. Each setting has one validation run. Validity stays zero. This is not a DiGress reproduction; gains over a small baseline need not stay the same at a larger budget.

4.6 Does augmentation help classification?

This secondary experiment asks whether generated graphs help a separate GNN predict SBM community count (2–5). “Real” graphs come from the SBM recipe. Augmentation adds one diffusion sample per real graph, using either no hint or high- $k=32$ hints. The three-layer dense Graph Isomorphism Network (GIN), hidden width 64, trains on 20/40/80 real graphs with five seeds. Communities have 20–40 nodes, $p_{\text{intra}} = 0.3$, $p_{\text{inter}} = 0.05$. We select models on 60 validation graphs and evaluate on 100 separate test graphs.

Neither diffusion method improves average accuracy over real-only training (Table 6). High beats no hints at 20 real graphs, but not at 40 or 80. Adding fresh SBM-recipe graphs instead reaches




Figure 2: Validation graph 0 and low- $k=8$ /unconditioned sample 0, seed 0. Panels share the reference’s node positions; this does not imply planarity. Solid green edges are shared; dashed gray edges are not. Labels show shared/generated counts and F_1 ; Table 5 covers all samples.

Real training graphs	Real only	+ No hint	+ High $k=32$
20	94.8	89.8	92.0
40	97.0	94.2	94.0
80	97.2	96.2	95.0

Table 6: Mean test accuracy (%), five seeds; higher is better. The last two columns add graphs from the unconditioned or high-frequency-conditioned generator. These are descriptive means, not significance tests.

98.8% at 80 real graphs, versus 97.2% for real-only.

This is a limited negative result, not evidence against augmentation in general. Generated graphs inherit template labels without checking their community count; label errors may hurt accuracy, but we did not measure their effect. Graph size is also an easy clue: a CPU nearest-class-mean classifier using only node count, fitted without tuning on the same training sets, scores 82.2/82.8/82.6% on the same test set. Unconditioned samples still receive node count, so inherited labels need not be independent of them. Finally, we tested high- $k=32$, not low-frequency augmentation, and did not check SBM samples for duplication. The experiment does not tell us whether correctly labeled, low-frequency-generated graphs would help.

5 Discussion and Limitations

A strength of this study is combining generation scores with recovery and copying checks. This helps distinguish improved statistics from reconstruction, although separating frequency from tar-

get information still needs further controls.

Generation is evaluated on Planar validation graphs. Selected settings have three seeds; most cutoffs have one. Generalization needs more seeds, datasets and untouched-test evaluation. The sweep supports aggregate comparisons; detailed component metrics are available for the pilot.

Our snapshot makes copying checks reproducible for conditioning graphs in fixed node order, not training graphs under arbitrary relabeling. Latent diffusion, cluster conditioning, combinatorial Laplacians and SignNet offer further comparisons. SBM results cover low- and high-band recovery, not a generation sweep; its shuffled control is excluded because graph sizes are not matched.

6 Conclusion

We studied spectral hints in graph diffusion and found repeatable improvements at low $k=8$ and high $k=32$ under oracle conditioning on Planar. Our main contribution is to pair generation scores with direct measurements of target recovery and source-graph similarity. These checks show why improvements should be read alongside reconstruction, rather than attributed to frequency alone. The code and sample snapshot make the copying analysis reusable. Together, the results and evaluation tools provide a practical starting point for testing spectral guidance across other graph families and model designs.

References

- Jacob Austin, Daniel D. Johnson, Jonathan Ho, Daniel Tarlow, and Rianne van den Berg. 2021. Structured denoising diffusion models in discrete state-spaces. In *Advances in Neural Information Processing Systems*.
- Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander Smola. 2012. A kernel two-sample test. *Journal of Machine Learning Research*, 13:723–773.
- Aric A. Hagberg, Daniel A. Schult, and Pieter J. Swart. 2008. Exploring network structure, dynamics, and function using NetworkX. In *Proceedings of the 7th Python in Science Conference*, pages 11–15.
- Jonathan Ho, Ajay Jain, and Pieter Abbeel. 2020. Denoising diffusion probabilistic models. In *Advances in Neural Information Processing Systems*, volume 33, pages 6840–6851.
- Jonathan Ho and Tim Salimans. 2022. Classifier-free diffusion guidance. *arXiv preprint arXiv:2207.12598*.
- Tomaž Hočevar and Janez Demšar. 2014. A combinatorial approach to graphlet counting. *Bioinformatics*, 30(4):559–565.
- Jaehyeong Jo, Seul Lee, and Sung Ju Hwang. 2022. Score-based generative modeling of graphs via the system of stochastic differential equations. In *Proceedings of the 39th International Conference on Machine Learning*, pages 10362–10383.
- Renjie Liao, Yujia Li, Yang Song, Shenlong Wang, William Hamilton, David K. Duvenaud, Raquel Urtasun, and Richard Zemel. 2019. Efficient graph generation with graph recurrent attention networks. In *Advances in Neural Information Processing Systems*, pages 4255–4265.
- Derek Lim, Joshua Robinson, Lingxiao Zhao, Tess E. Smidt, Suvrit Sra, Haggai Maron, and Stefanie Jegelka. 2023. Sign and basis invariant networks for spectral graph representation learning. In *International Conference on Learning Representations*.
- Karolis Martinkus, Andreas Loukas, Nathanaël Perraudin, and Roger Wattenhofer. 2022. SPECTRE: Spectral conditioning helps to overcome the expressivity limits of one-shot graph generators. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*.
- Giorgia Minello, Alessandro Bicciato, Luca Rossi, Andrea Torsello, and Luca Cosmo. 2025. Generating graphs via spectral diffusion. In *The Thirteenth International Conference on Learning Representations*.
- Alexander Quinn Nichol and Prafulla Dhariwal. 2021. Improved denoising diffusion probabilistic models. In *Proceedings of the 38th International Conference on Machine Learning*, pages 8162–8171.
- Chenhao Niu, Yang Song, Jiaming Song, Shengjia Zhao, Aditya Grover, and Stefano Ermon. 2020. Permutation invariant graph generation via score-based generative modeling. In *Proceedings of the 23rd International Conference on Artificial Intelligence and Statistics*, volume 108, pages 4474–4484.
- Clement Vignac, Igor Krawczuk, Antoine Siraudin, Bohan Wang, Volkan Cevher, and Pascal Frossard. 2023. DiGress: Discrete denoising diffusion for graph generation. In *International Conference on Learning Representations*.
- Yuhang Xie and Sinno Jialin Pan. 2026. [Global and local topology-aware graph generation via dual conditioning diffusion](#). In *International Conference on Learning Representations*.
- Jiaxuan You, Rex Ying, Xiang Ren, William L. Hamilton, and Jure Leskovec. 2018. GraphRNN: Generating realistic graphs with deep auto-regressive models. In *Proceedings of the 35th International Conference on Machine Learning*.
- Cai Zhou, Xiyuan Wang, and Muhan Zhang. 2024. Unifying generation and prediction on graphs with latent graph diffusion. In *Advances in Neural Information Processing Systems*.
- Junyou Zhu, Langzhou He, Chao Gao, Dongpeng Hou, Zhen Su, Philip S. Yu, Jürgen Kurths, and Frank Hellmann. 2025. SDMG: Smoothing your diffusion models for powerful graph representation learning. In *Proceedings of the 42nd International Conference on Machine Learning*.